

Biostat 276 Homework 1

Zian ZHUANG

Bayesian Probit Regression

In R load the package (survival) and consider the analysis of the data-set (infert). Ignoring dependence due to matching, consider a Bayesian analysis for a logistic regression model relating case status to: age, parity, education, spontaneous and induced. More precisely, assume case status y_i has density $y_i \sim_{ind} Bern(p_i), p_i = \Phi(X_i' \beta)$, where $\Phi(\cdot)$ is the standard Gaussian cdf. Consider a prior $\beta \sim N(0, 10^2(X'X)^{-1})$. We are interested in $p(\beta|Y)$.

```
library(survival)
data("infert")
infert$education <- infert$education %>% as.numeric() - 1
df <- infert %>%
  select(c("case", "education", "age", "parity", "induced", "spontaneous"))
```

(1)

Describe and implement an adaptive Metropolis-Hastings algorithm designed to obtain a MC with stationary distribution $p(\beta|Y)$.

We know that

$$\begin{aligned} P(\beta|Y) &= P(Y|\beta)P(\beta) \\ &= \left(\prod_{i=1}^n \Phi(X_i' \beta)^{y_i} (1 - \Phi(X_i' \beta))^{1-y_i} \right) P(\beta) \\ &\propto \left(\prod_{i=1}^n \Phi(X_i' \beta)^{y_i} (1 - \Phi(X_i' \beta))^{1-y_i} \right) * \exp\left(-\frac{1}{2} \beta' \sum \beta\right) \end{aligned}$$

Then we design a adaptive proposal:

- $\beta_{t+1}|\beta_t \sim N(\beta_t, c * \sum_t)$

in which $\sum_t = \frac{1}{t} \sum_{j=1}^t \beta_j \beta_j'$.

According to (Roberts and Rosenthal, 2009), we also set,

- 1) $\tilde{\sum}_t = \sum_t + \epsilon I$ to avoid degeneracies
- 2) $Q(\beta_t, \beta_{t+1}) = \delta N(\beta_t, c * \tilde{\sum}_t) + (1 - \delta) N(\beta_t, \sum_0)$ in which $\delta \sim Binomial(p_\delta)$. here we set $p_\delta = 0.7$, and set \sum_0 same as the prior \sum of β .

```

bvar_prior <- solve(t(as.matrix(df[, -1])) %*% as.matrix(df[, -1])) * 100
bvar_prior_inv <- t(as.matrix(df[, -1])) %*% as.matrix(df[, -1]) / 100
.cal_ll <- function(beta){
  est <- data.frame(est=pnorm(as.matrix(df[, -1]) %*% t(beta)), y=df[, 1])
  out <- apply(est, 1, function(x){ifelse(x[2]==0, 1-x[1], x[1])})
  out <- log(out) %>% sum - 1/2 * beta %*% bvar_prior_inv %*% t(beta)
}

mh.sim <- function(nsim=1000, burn=0, delta=0.8, c=3, seed=1996){
  set.seed(seed)

  num_beta <- dim(df)[2]-1
  nsim.total <- nsim*(1.0 + burn)
  burn.num <- nsim*burn
  delta_set <- rbinom(nsim.total, 1, prob = delta)

  beta <- matrix(rep(0, num_beta), 1)
  beta.ch <- matrix(NA, nsim, num_beta)
  bvar <- diag(num_beta)*10^(-15)

  for(i in 1:nsim.total){
    if(i <= 1500){ #used adaptive algorithm on the first 1500 simulations
      bvar <- (bvar * (i - 1) + t(beta) %*% beta)/i + diag(num_beta)*10^(-15)
      beta_temp_2 <- rmvnorm(n=1, mean = beta, sigma = bvar_prior)
      delta_temp <- delta_set[i]
    }else{
      delta_temp <- 1
    }
    beta_temp_1 <- rmvnorm(n=1, mean = beta, sigma = c*bvar)
    beta_temp <- beta_temp_1 * delta_temp + beta_temp_2 * (1-delta_temp)
    P0 <- .cal_ll(beta)
    P1 <- .cal_ll(beta_temp)
    ratio <- P1 - P0
    if(log(runif(1)) < ratio){
      beta <- beta_temp
    }
    if(i > burn.num){
      i1 <- i - burn.num
      beta.ch[i1,] = beta
    }
  }
  return(list(beta=beta.ch))
}

```

simulation begin

```

burn=0.3
nsim=5000

# simulation
mh.mcmc.sim <- mh.sim(nsim=nsim, burn=burn, delta=0.7, c=1)

# get results

```

```
results.mh <- apply(mh.mcmc.sim$beta, 2,
  function(x){
    quantile(x, c(0.025, 0.5, 0.975))) %>% round(.,2) %>%
    apply(., 2, function(x){paste0(x[2], " (", x[1], ", ", x[3], ")")})})
```

(2)

Describe and implement a data augmented (DA-MCMC) strategy targeting $p(\beta|Y)$.

We add a new parameter z here, which makes

$$Y_i|z_i \sim I(z_i > 0)$$

$$z_i|\beta, \tau_i \sim N(X_i\beta, 1)$$

Then we can calculate the full posterior distribution for β ,

$$\begin{aligned} P(\beta|z, \tau) &= P(z|\beta) * P(\beta) \\ &\propto \prod \exp\left(-\frac{(z_i - X_i\beta)^2}{2}\right) * \exp\left(-\frac{1}{2}\beta' \sum \beta\right) \\ &= \exp\left(-\frac{1}{2}(\beta' X' X \beta - 2 * \beta X' z) - \frac{1}{2} * \beta' \sum \beta\right) \\ &= \exp\left(-\frac{1}{2}(\beta' (X' X + \sum) \beta - 2 * \beta X' z)\right) \end{aligned}$$

This is a normal kernel. Then we found that $\beta|z \sim N\left(\frac{X'z}{X'X+\sum}, \frac{1}{X'X+\sum}\right)$.

As for z ,

$$\begin{aligned} P(z_i|\beta, Y) &= P(Y_i|z_i) * P(z_i|\beta) \\ &= I(z_i > 0) * \exp\left(-\frac{(z_i - X_i\beta)^2}{2}\right) \text{ (if } Y_i=1) \\ &= I(z_i \leq 0) * \exp\left(-\frac{(z_i - X_i\beta)^2}{2}\right) \text{ (if } Y_i=0) \end{aligned}$$

Then we found that z_i follows truncated normal distribution with mean $X_i\beta$ and variance 1.

```
da.mh.sim <- function(nsim=1000, burn=0, seed=1996,
  X=as.matrix(df[, -1]),
  Y=as.matrix(df[, 1])){
  set.seed(seed)

  num_beta <- dim(df)[2]-1
  nsim.total <- nsim*(1.0 + burn)
  burn.num <- nsim*burn

  beta <- matrix(rep(0, num_beta), 1)
  beta.ch <- matrix(NA, nsim, num_beta)
  z <- rep(0.5, length(Y))

  range <- cbind(ifelse(Y==0, -Inf, 0),
    ifelse(Y==0, 0, Inf))
```

```

for(i in 1:nsim.total){
  #i=1
  #z
  z_info <- as.matrix(cbind(X%*%t(beta), # z_mean
                           1, # z_sd
                           range),) # truncated range
  z <- apply(z_info, 1,
            function(info){rtruncnorm(1, a=info[3], b=info[4],
                                       mean = info[1], sd = info[2])})

  #beta
  bmean <- solve(t(X)%*%X+bvar_prior)%*%t(X)%*%z
  bvar <- solve(t(X)%*%X+bvar_prior)
  beta <- rmvnorm(n=1, mean = bmean, sigma = bvar)

  if(i > burn.num){
    i1 <- i - burn.num
    beta.ch[i1,] = beta
  }
}
return(list(beta=beta.ch))
}

```

simulation begin

```

burn=0.3
nsim=5000

# simulation
da.mh.mcmc.sim <- da.mh.sim(nsim=nsim, burn=burn)

# get results
results.da <- apply(da.mh.mcmc.sim$beta, 2,
                  function(x){
                    quantile(x, c(0.025, 0.5, 0.975))) %>% round(.,2) %>%
                    apply(., 2, function(x){paste0(x[2], " (", x[1], ", ", x[3], ")")})

```

(3)

Describe and implement a parameter expanded - data augmentation (PX-DA MCMC) algorithm targeting $p(\beta|Y)$.

Here we add w and α , which follow

$$\begin{aligned}
 W_i | \beta, \alpha &\sim N(X_i' \beta \alpha, \alpha^2) \\
 Y_i | W_i &= I(W_i > 0) \\
 \alpha^2 &\sim IG(a, b)
 \end{aligned}$$

Then we can calculate the full posterior distribution for β ,

$$\begin{aligned}
P(\beta|w, \alpha) &= P(w|\beta, \alpha) * P(\beta) \\
&\propto \prod \exp(-\frac{(w_i - \alpha X_i \beta)^2}{2\alpha^2}) * \exp(-\frac{1}{2} \beta' \sum \beta) \\
&= \exp(-\frac{1}{2\alpha^2} (\alpha^2 \beta' * X' X \beta - 2 * \alpha \beta X' w) - \frac{1}{2} * \beta' \sum \beta) \\
&= \exp(-\frac{1}{2} (\beta' (X' X + \sum) \beta - 2 * \beta X' w / \alpha))
\end{aligned}$$

This is a normal kernel. Then we found that $\beta|z, \tau \sim N(\frac{X' w}{\alpha X' X + \alpha \sum}, \frac{1}{X' X + \sum})$.

As for w ,

$$\begin{aligned}
P(w_i|\beta, \alpha, Y_i) &= P(Y_i|w_i) * P(w_i|\beta, \alpha) \\
&= I(w_i > 0) * \exp(-\frac{(w_i - \alpha X_i \beta)^2}{2\alpha^2}) \text{ (if } Y_i=1) \\
&= I(w_i \leq 0) * \exp(-\frac{(w_i - \alpha X_i \beta)^2}{2\alpha^2}) \text{ (if } Y_i=0)
\end{aligned}$$

Then we found that w_i follows truncated normal distribution with mean $\alpha X_i \beta$ and variance α^2 .

As for α^2 ,

$$\begin{aligned}
P(\alpha^2|\beta, w, Y) &= P(w|\alpha^2, \beta) * P(\alpha^2) \\
&= (\frac{1}{\sqrt{2\pi\alpha^2}})^n \exp(-\frac{(w - \alpha X \beta)^2}{2\alpha^2}) * (\alpha^2)^{-a-1} \exp(-b/\alpha^2) \\
&\propto (\alpha^2)^{-a-n/2-1} \exp(-(\frac{(w - \alpha X \beta)^2}{2} + b)/\alpha^2)
\end{aligned}$$

This is difficult to treat it as a specific kernel. Then we consider use MH method to simulate α^2 . Here we use truncated normal proposal.

```

.alpha.sq.ll <- function(alpha.sq,beta,w,a,b,X,n){
  out <- (-a-n/2-1)*log(alpha.sq)-
    ((2*b+t(w-sqrt(alpha.sq)*X%%t(beta))%%
      (w-sqrt(alpha.sq)*X%%t(beta)))/2)/alpha.sq
  return(out)
}
px.da.mh.sim <- function(nsim=1000, burn=0, as_sd=1,
  a=3, b=3, seed=1996,
  X=as.matrix(df[, -1]),
  Y=as.matrix(df[, 1])){
  set.seed(seed)

  num_beta <- dim(df)[2]-1
  nsim.total <- nsim*(1.0 + burn)
  burn.num <- nsim*burn

  beta <- matrix(rep(0,num_beta),1)
  beta.ch <- matrix(NA,nsim,num_beta)
  alpha.sq <- 1
  n <- length(Y)

```

```

range <- cbind(ifelse(Y==0, -Inf, 0),
               ifelse(Y==0, 0, Inf))

for(i in 1:nsim.total){
  #i=1
  #w
  w_info <- as.matrix(cbind(sqrt(alpha.sq)*X%%t(beta), # w_mean
                            sqrt(rep(alpha.sq,length(Y))), # w_sd
                            range),) # truncated range

  w <- apply(w_info, 1,
            function(info){rtruncnorm(1, a=info[3], b=info[4],
                                     mean = info[1], sd = info[2])})

  #beta
  bmean <- solve(t(X)%*%X*sqrt(alpha.sq)+
                bvar_prior*sqrt(alpha.sq))%*%t(X)%*%w
  bvar <- solve(t(X)%*%X+bvar_prior)
  beta <- rmvnorm(n=1, mean = bmean, sigma = bvar)

  #alpha.sq
  as_temp <- rtruncnorm(1, a=0, b=Inf, mean = alpha.sq, sd = as_sd)
  P0 <- .alpha.sq.ll(alpha.sq,beta,w,a,b,X,n)
  P1 <- .alpha.sq.ll(as_temp,beta,w,a,b,X,n)
  ratio <- P1 - P0 - log(dtruncnorm(as_temp,mean = alpha.sq, sd = as_sd)) +
    log(dtruncnorm(alpha.sq,mean = as_temp, sd = as_sd))

  if(log(runif(1)) < ratio){
    alpha.sq <- as_temp
  }

  if(i > burn.num){
    i1 <- i - burn.num
    beta.ch[i1,] <- beta
  }
}
return(list(beta=beta.ch))
}

```

simulation begin (set as_sd=1, a=3, b=3)

```

burn=0.3
nsim=5000

# simulation
px.da.mh.mcmc.sim <- px.da.mh.sim(nsim=nsim, burn=burn,as_sd=1)

# get results
results.px <- apply(px.da.mh.mcmc.sim$beta, 2,
                  function(x){
                    quantile(x, c(0.025, 0.5, 0.975))) %>% round(.,2) %>%
                    apply(., 2, function(x){paste0(x[2]," (", x[1],", ",x[3],")")})

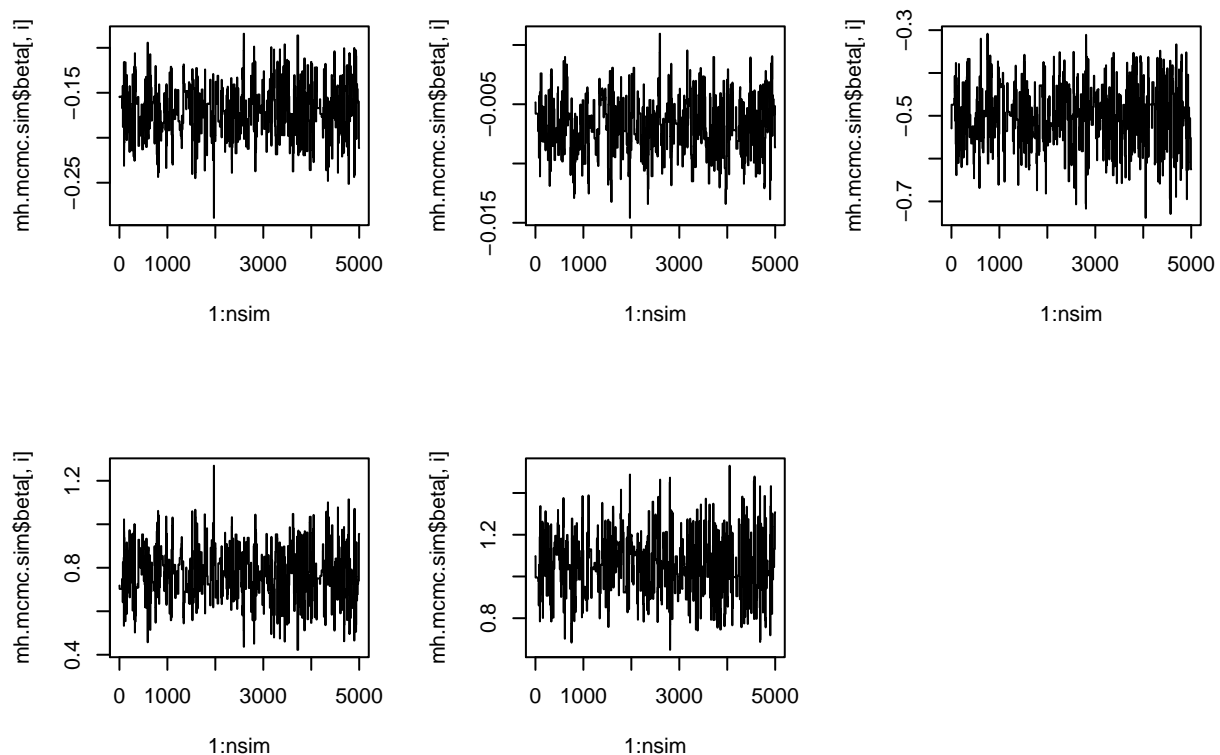
```

(4)

Assess mixing and convergence of the chains induced by the competing transition schemes implemented in 1,2 and 3. Comment on potential trade-offs involving: coding complexity, storage and cpu time.

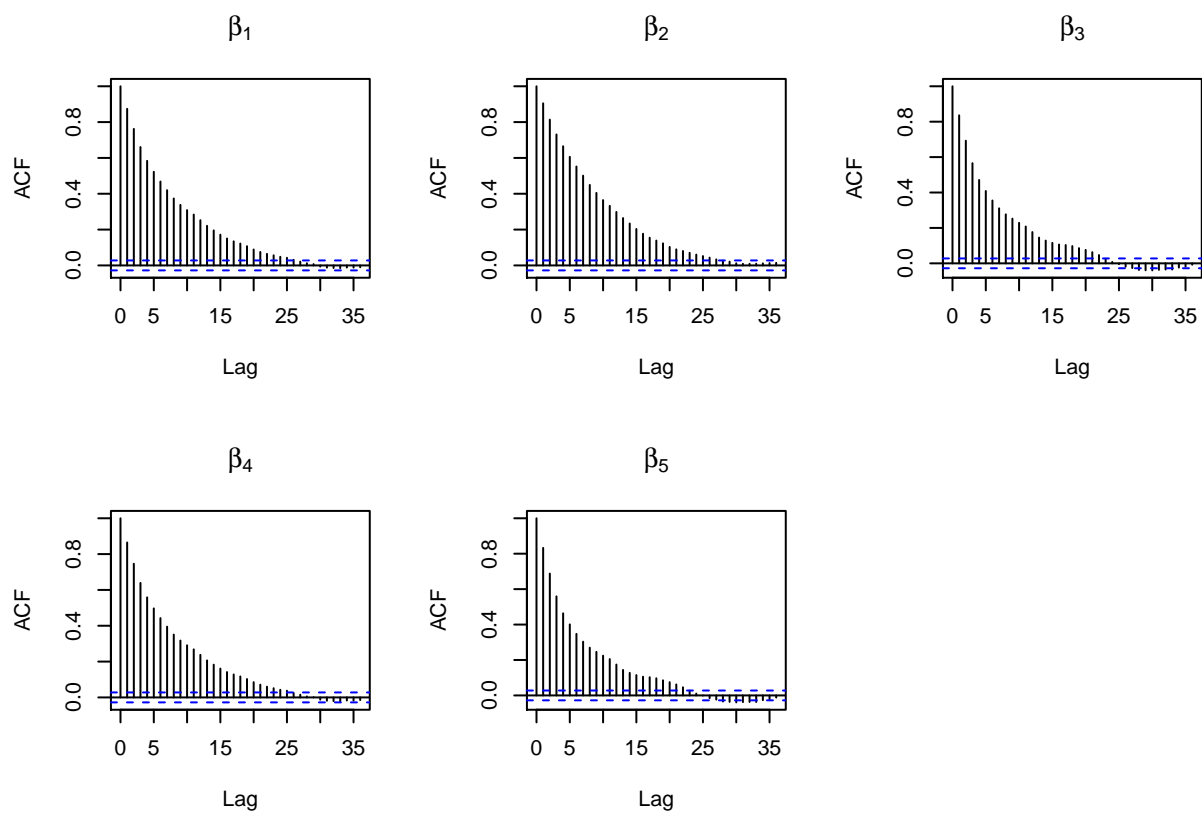
Check mixing and convergence of the chains:

```
# plots for mh method
par(mfrow=c(2,3))
for(i in 1:5){
  plot(1:nsim, mh.mcmc.sim$beta[,i], type="l")
}
par(mfrow=c(2,3))
```

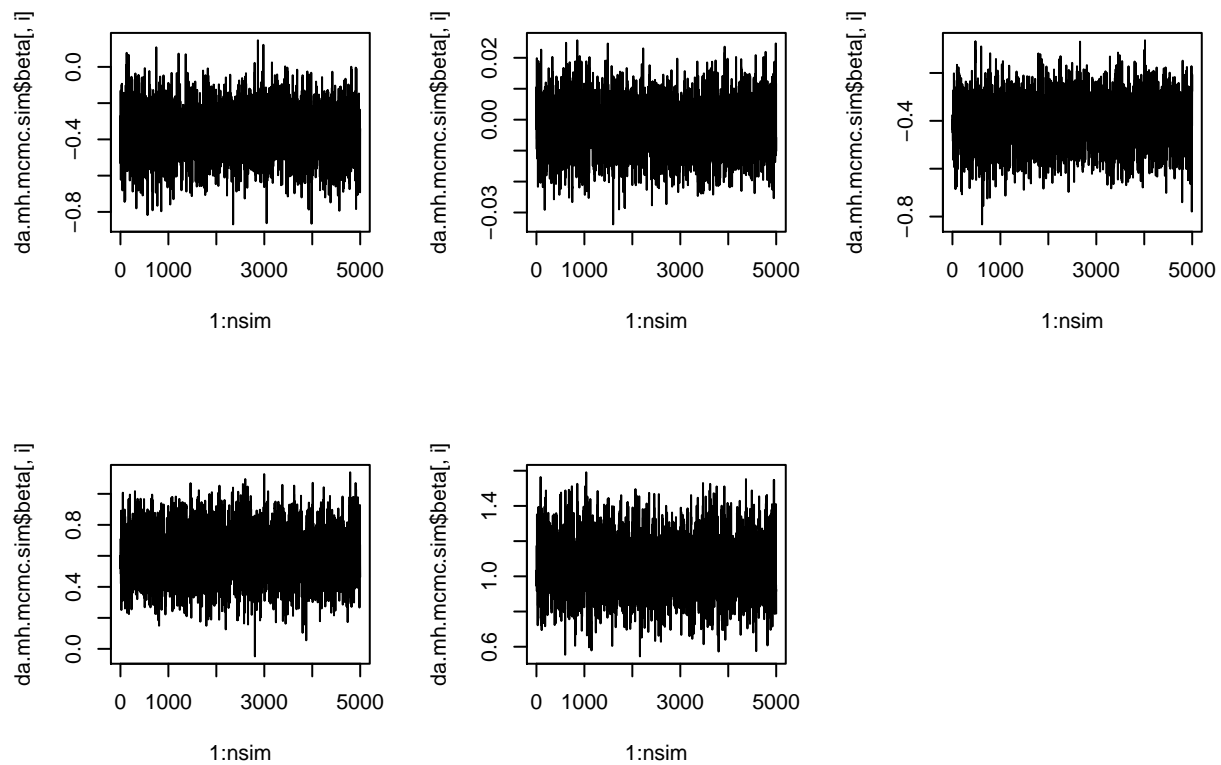


```
for(i in 1:5){
  acf(mh.mcmc.sim$beta[,i], main = substitute(beta[x], list(x = i)))
}

# plots for da-mh method
par(mfrow=c(2,3))
```

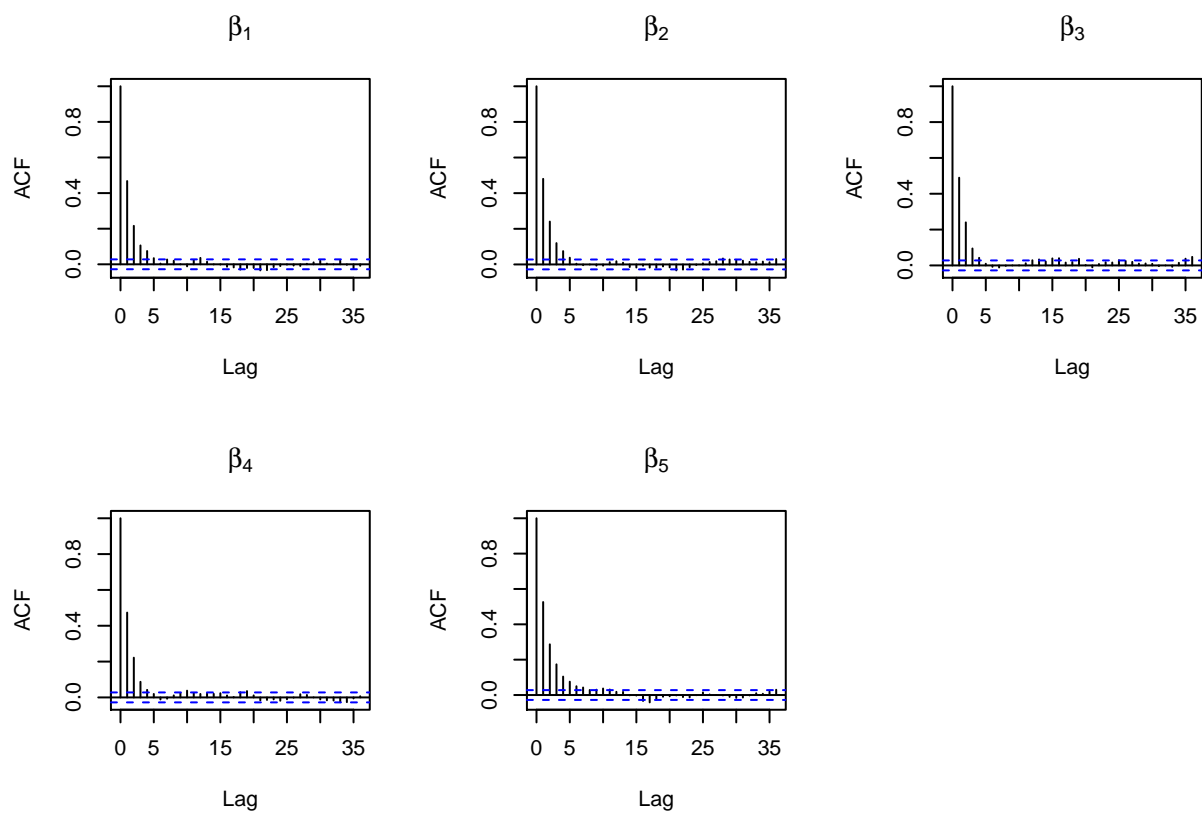


```
for(i in 1:5){
  plot(1:nsim, da.mh.mcmc.sim$beta[,i], type="l")
}
par(mfrow=c(2,3))
```

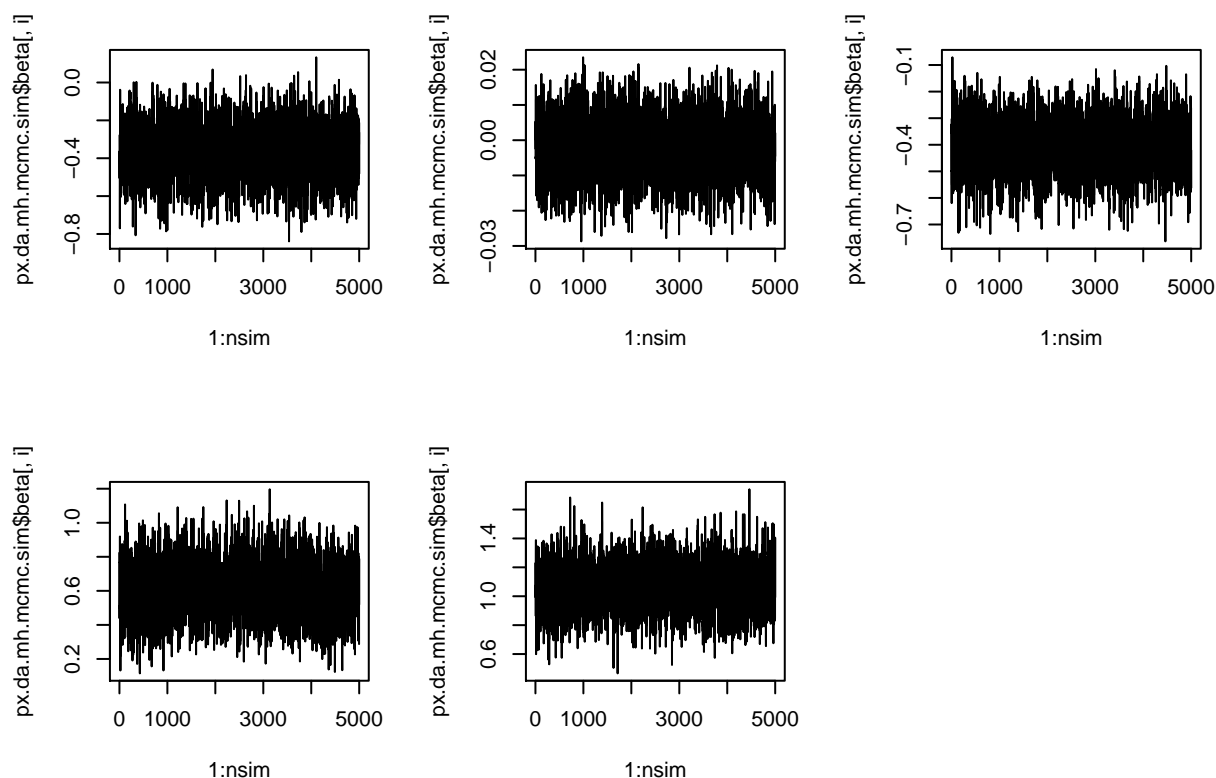



```
for(i in 1:5){
  acf(da.mh.mcmc.sim$beta[,i], main = substitute(beta[x], list(x = i)))
}

# plots for px-da-mh method
par(mfrow=c(2,3))
```



```
for(i in 1:5){
  plot(1:nsim, px.da.mh.mcmc.sim$beta[,i], type = "l")
}
par(mfrow=c(2,3))
```



```
for(i in 1:5){
  acf(px.da.mh.mcmc.sim$beta[,i], main = substitute(beta[x], list(x = i)))
}

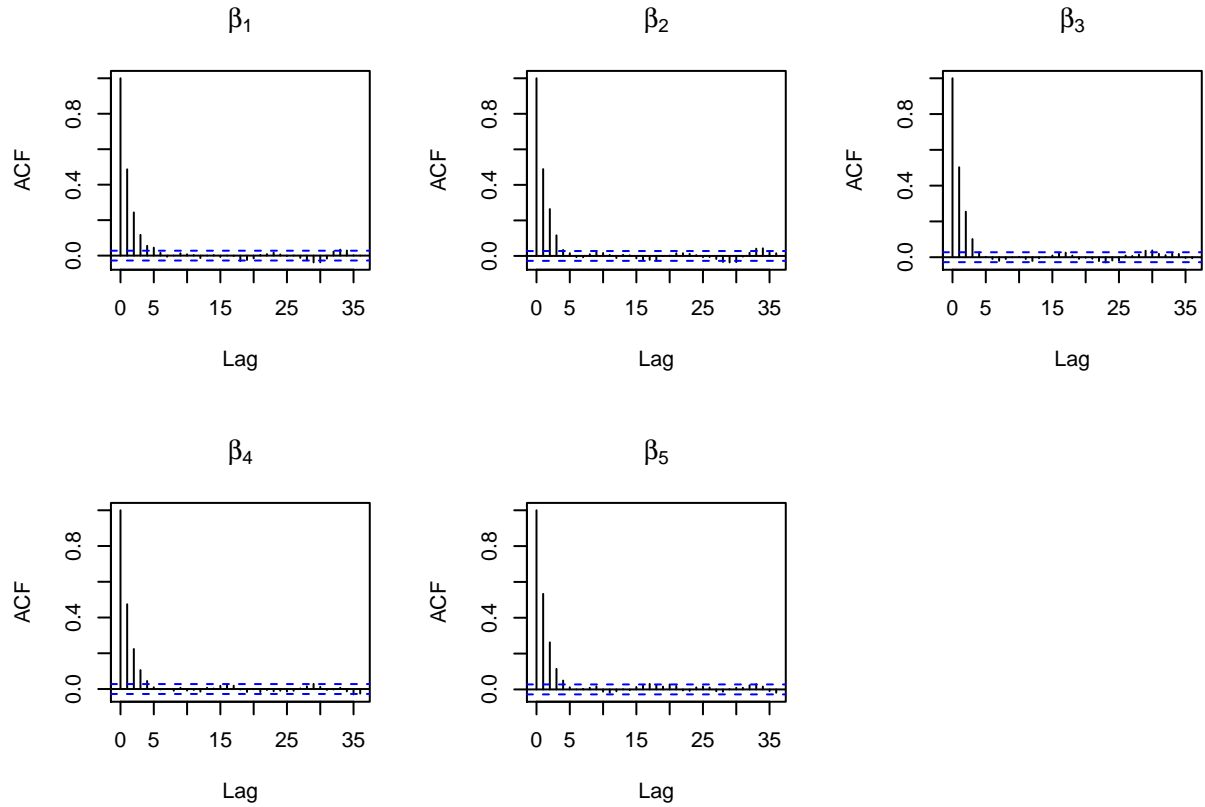
results <- rbind(results.mh, results.da, results.px) %>% as.data.frame
colnames(results) <- colnames(df)[-1]
results %>%
  kbl(caption = "Beta Summary Table (95% CI)") %>%
  kable_classic(full_width = F, html_font = "Cambria")
```

\begin{table}

\caption{Beta Summary Table (95% CI)}

	education	age	parity	induced	spontaneous
results.mh	-0.17 (-0.23, -0.12)	-0.01 (-0.01, 0)	-0.5 (-0.64, -0.37)	0.76 (0.55, 1.02)	1.05 (0.79, 1.32)
results.da	-0.36 (-0.63, -0.1)	0 (-0.02, 0.01)	-0.41 (-0.6, -0.22)	0.61 (0.31, 0.92)	1.06 (0.77, 1.37)
results.px	-0.37 (-0.64, -0.11)	0 (-0.02, 0.01)	-0.4 (-0.6, -0.22)	0.61 (0.31, 0.92)	1.06 (0.75, 1.38)

\end{table}



As we can tell from the plots, auto correlation of chain generated by `adaptive-mh` method decrease much slower than that of `px-da-mh` and `da-mh` method. In addition, auto correlation for some parameters' chains (e.g. β_5) of `px-da-mh` decrease slightly faster than that of `da-mh` method. As for values of β s, three methods provided largely consistent estimates.

Calculate memory usage and running time of the chain:

```
# set same situation
burn=0.3
nsim=500

# calculate memory usage
mem.mh <- profmem(mh.sim(nsim=nsim, burn=burn))
mem.da.mh <- profmem(da.mh.sim(nsim=nsim, burn=burn))
mem.px.da.mh <- profmem(px.da.mh.sim(nsim=nsim, burn=burn))

mem_use <- data.frame(method=c("adaptive-mh", "da.mh", "px.da.mh"),
                      memory_usage=c(mem.mh$bytes %>% sum(na.rm = T),
                                      mem.da.mh$bytes %>% sum(na.rm = T),
                                      mem.px.da.mh$bytes %>% sum(na.rm = T)))

mem_use %>%
  kbl(caption = "Memory Usage Summary Table") %>%
  kable_classic(full_width = F, html_font = "Cambria")
```

Table 1: Memory Usage Summary Table

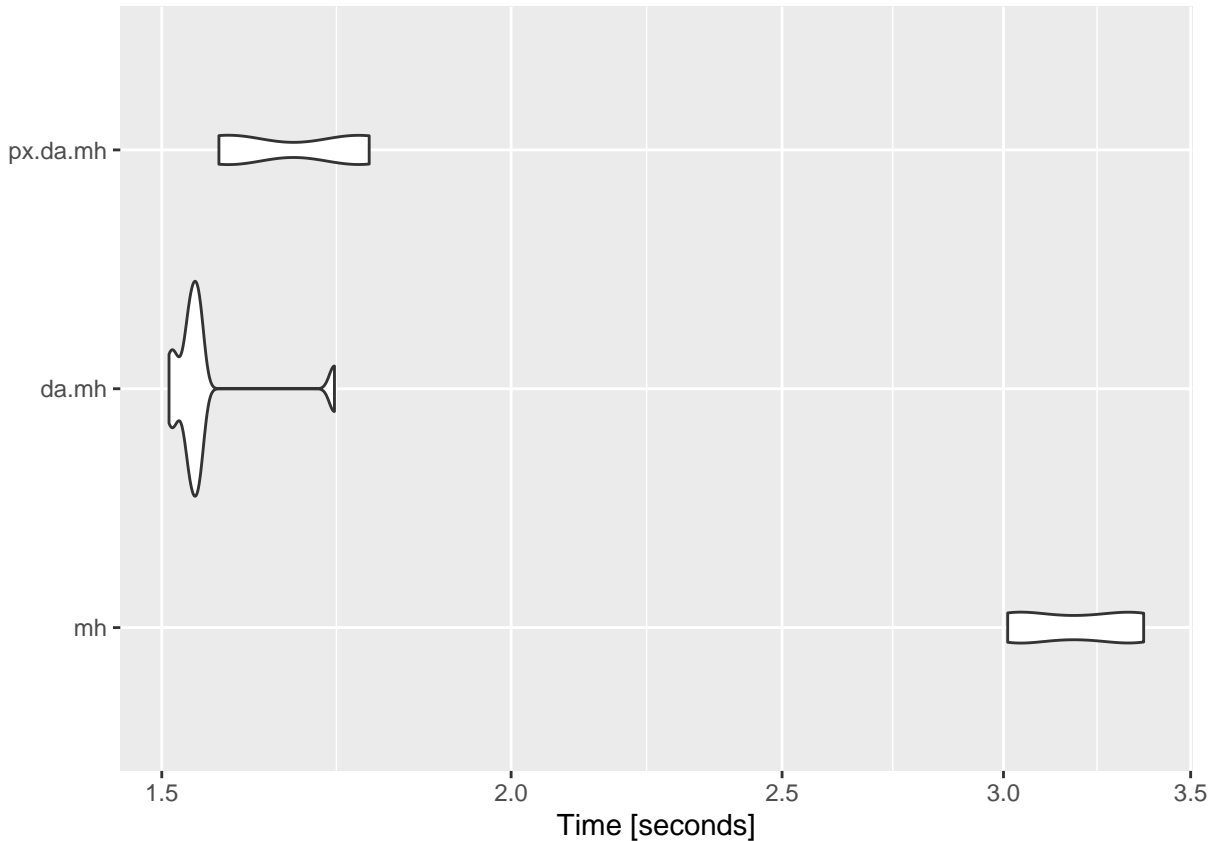
method	memory_usage
adaptive-mh	58159376
da.mh	460175672
px.da.mh	472736840

```
# calculate running time
stats <- microbenchmark(mh = {mh.sim(nsim=nsim, burn=burn)},
                        da.mh = {da.mh.sim(nsim=nsim, burn=burn)},
                        px.da.mh = {px.da.mh.sim(nsim=nsim, burn=burn)},
                        times = 10, unit = "ms")
stats
```

```
## Unit: milliseconds
##      expr      min       lq     mean  median      uq      max neval cld
##      mh 3010.472 3028.619 3185.825 3184.099 3341.571 3367.114   10   b
##     da.mh 1508.983 1530.647 1554.086 1538.988 1548.044 1729.210   10   a
##    px.da.mh 1572.284 1582.896 1674.338 1671.075 1763.771 1779.397   10   a
```

```
autoplot(stats)
```

Coordinate system already present. Adding new coordinate system, which will replace the existing one



As we can tell from the memory usage summary table, **adaptive-mh** method used the smallest memory. **da.mh** and **px.da.mh** methods used very similar memory. As for the efficiency, **da.mh** and **px.da.mh** have similar running time, which is nearly half of **adaptive-mh** method's.

Conclusion: **px.da.mh** method had the best convergence of the chains. **da.mh** method operated most efficient and **adaptive-mh** method occupied the smallest memory.